

AMAN NIKUMB

Software Engineer — Data

Gurugram, India | +91 85140 87234 | amannikumbh73@gmail.com | linkedin.com/in/aman-nikumb | github.com/aman0981

Backend & data engineer with 1.5+ years building scalable APIs, distributed systems, and code-first ETL. Strong in Python (FastAPI, Django, Celery), PySpark/Databricks medallion pipelines, PostgreSQL/Elasticsearch and Docker. Promoted to Software Engineer — Data (Apr 2026). Aspiring Forward Deployed Engineer: requirements → build & ship → deploy → data → business analytics.

EXPERIENCE

Software Engineer — Data, Engineo Solutions

Apr 2026 — Present

- Build a code-first ETL pipeline on Databricks with PySpark, transforming data through a bronze → silver → gold medallion architecture and loading the curated output into Elasticsearch.
- Expose a streaming API that delivers data to clients in chunks, with Elasticsearch queries built dynamically from each client request's filter payload.
- Scaling toward a centralized API integrating an external data endpoint, a parse service to normalize responses, and a DocumentDB sink — delivering to clients while persisting for reuse.

Associate Software Engineer — Data, Engineo Solutions

Nov 2024 — Apr 2026

- Built a distributed data-extraction platform for a US travel-domain client (Django, FastAPI, Celery, RabbitMQ, Redis, MongoDB, PostgreSQL) for scalable aggregation and monitoring.
- Engineered an automated ingestion pipeline for a Big-4 consulting firm — Excel → JSON → PostgreSQL via APIs — cutting manual effort ~70%, with parent-child relational mapping and FK resolution.
- Implemented Redis caching and rate-limiting (40%+ fewer redundant calls) and a layered API-security stack (API-key, IP allow-listing, JWT, CORS, bad-user-agent flagging); integrated AWS S3 + OpenSearch.

PROJECTS

XBRL Intelligence Engine — Solo

github.com/aman0981/XBRL_Engine

Python · FastAPI · async SQLAlchemy 2.0 / asynccpg · PostgreSQL 16 · Alembic · structlog · Docker

- Async, multi-source ETL ingesting XBRL/iXBRL from SEC EDGAR, EU ESEF and direct upload into one canonical financial model — 24,852 facts, 71 filings, 552 canonical line-items.
- Mode A vs Mode B cross-source reconciliation + bitemporal restatement detection (dimension/unit/decimals-aware); 135 automated tests; hardened non-root multi-stage Docker image.

HireBeacon — India Job-Intelligence Platform — Solo

github.com/aman0981/HireBeacon

Python · Supabase (Postgres 17, pgvector) · Next.js 15 · TypeScript · GitHub Actions · Groq/Gemini

- Unified 13 ATS sources (incl. 4 of 5 enterprise "big-5") behind one adapter contract; idempotent, self-healing ingestion that never deletes live inventory on a partial crawl.
- Semantic résumé→job matching: MiniLM 384-dim embeddings → pgvector cosine top-25 via a Postgres RPC → fallback-chained LLM rationale; row-level security; 336 tests; free-tier infra.

TECHNICAL SKILLS

Languages: Python, SQL, Golang

Backend & APIs: FastAPI, Django, DRF, Celery, Pydantic, REST, JWT / API-key security

Data & ETL: PySpark, Databricks, Medallion ETL, Pandas, web crawling (Playwright, BeautifulSoup)

Datastores & Search: PostgreSQL, MongoDB, Redis, Elasticsearch / OpenSearch, pgvector, DocumentDB, AWS S3

Async & Infra: asyncio, async SQLAlchemy 2.0, RabbitMQ, Docker, AWS, GitHub Actions, Git

EDUCATION

B.Tech, Computer Science

2020 — 2024

GITA Autonomous College · GPA 8.54

ACHIEVEMENTS

200+ problems solved on LeetCode & HackerRank · LeetCode 50-Days & Top SQL 50 badges.

Certs: Become a Django Developer (LinkedIn) · 100 Days of Code: Python (Udemy)